

emcee: Ensemble MCMC for the masses

Daniel Foreman-Mackey¹, David W. Hogg, *et al.*

*Center for Cosmology and Particle Physics, Department of Physics, New York University,
4 Washington Place, New York, NY, 10003, USA*

ABSTRACT

emcee is a stable, well tested Python implementation of the affine-invariant ensemble sampler for Markov chain Monte Carlo (MCMC) proposed by Goodman & Weare (2010). The code is open source and has already been used in several published projects in the Astrophysics literature. The algorithm behind **emcee** has several advantages over traditional MCMC sampling methods and it has excellent performance measured by the sampling autocorrelation time on several test problems. One major advantage of the algorithm is that it requires the hand-tuning of only 2 parameters compared to the $\sim N^2$ for a traditional algorithm in N dimensions. In this *Article*, we describe the Goodman & Weare algorithm and the details of our implementation. **emcee** takes advantage of the naturally parallel nature of the algorithm allowing *any* user to take advantage of multiple CPUs without any extra effort. We discuss the subtleties associated with this parallelization.

The code is available online at <http://danfm.ca/emcee> under the GNU Public License (v2).

Subject headings: methods: data analysis — methods: numerical — methods: statistical

1. Introduction

Probabilistic data analysis is a common theme in the recent scientific literature since. In particular, many problems in astrophysics and cosmology, benefit from careful probabilistic inference because most problems of interest lie in the regime of very low signal-to-noise data with large systematic uncertainties and missing observations. In many scientific problems of interest, however, there is a physically motivated, generative model of the data or several

¹danfm@nyu.edu

such models that we wish to compare. If we can leverage the physics that we understand and the rich prior information from the literature then there is a chance — it has been repeatedly shown — of solving these problems. With the use of more sophisticated models for the data comes an engineering challenge that can only be solved with a combination of computing power and more efficient algorithms.

The general probabilistic data analysis procedure involves examining either the posterior probability distribution function (PDF) of the parameters of the model or the likelihood function of the data. In some cases, it is sufficient to just find the maximum of this function but it is often of interest to (at least approximately) understand the details of the function. In particular, if we wish to propagate the effects of measurement uncertainties through to the result of the analysis, we must sample the PDF on sufficiently small scales. Markov chain Monte Carlo (MCMC) methods are designed to approximate this function even in high dimensional parameter spaces and the use of MCMC has proved useful in too many research applications to list here (e.g. WMAP, etc. DFM: add citations).

Arguably the most important advantage of fully probabilistic data analysis is that it is possible to *marginalize* over nuisance parameters. A nuisance parameter is a component of the generative model that is of little or physical interest but must be specified in order to generate the data. Marginalization is the process of integrating over all possible values of these parameters and hence propagating the effects of your uncertainty about the true value to the final result. This operation is written as the integral

$$p(\mathbf{X}|\Theta) = \int p(\mathbf{X}|\Theta, \alpha) p(\alpha) d\alpha \quad (1)$$

where $p(\mathbf{X}|\Theta)$ is the marginalized likelihood of the data \mathbf{X} given the model parameters Θ and α is vector of nuisance parameters. In some specific cases, the integral in Equation (1) is analytically tractable but in general, the likelihood function $p(\mathbf{X}|\Theta, \alpha)$ is not a simple integrable function. In fact, in many problems, the likelihood function is actually the result of an extremely expensive numerical simulation. In this regime, the integration must be calculated numerically and it is often relatively efficient to approximate Equation (1) using MCMC sampling. If the likelihood function is expensive to calculate, it is advantageous to use a sampling algorithm that reduces the necessary number of likelihood evaluations. This also precludes the use of second order methods (such as hybrid/Hamiltonian Monte Carlo) that require the calculation of (numerical) gradients of the likelihood function.

Most uses of MCMC in the astrophysics literature are based on slight modifications to the Metropolis-Hastings (M-H) method (e.g. [MacKay 2003](#)). Each step in a M-H chain is proposed using a multivariate Gaussian centered on the current position of the chain. Since each term in the covariance matrix of this proposal distribution is an unspecified parameter,

this method has $D[D+1]/2$ tuning parameters (where D is the dimension of the parameter space). To make matters worse, the performance of this sampler is very sensitive to the optimality of these tuning parameters and there is no fool-proof method for choosing the values correctly. As a result, many heuristic methods have been developed to attempt to determine the optimal parameters in a data-driven way (e.g. Gregory 2005; Dunkley *et al.* 2005; Widrow *et al.* 2008). Unfortunately, these methods all require “burn-in” phases where shorter Markov chains are sampled and the results are used to tune the hyperparameters. This extra cost is unacceptable when the likelihood calls are computationally heavy.

The problem with traditional sampling methods can be visualized by studying the highly anisotropic density

$$p(\mathbf{x}) \propto \exp\left(-\frac{(x_1 - x_2)^2}{2\epsilon} - \frac{(x_1 + x_2)^2}{2}\right) \quad (2)$$

which would be considered “difficult” by standard MCMC algorithms. Equation (2) can be transformed into the much easier problem of sampling an isotropic Gaussian by an *affine transformation* of the form

$$y_1 = \frac{x_1 - x_2}{\sqrt{\epsilon}}, \quad y_2 = x_1 + x_2. \quad (3)$$

Therefore, an algorithm that is *affine invariant* will be insensitive to covariances between parameters. An affine invariant algorithm is unaffected by any transformation of the density of the form

$$\mathbf{Y} = \mathbf{A} \mathbf{X} + \mathbf{b}. \quad (4)$$

Extending earlier work by Christen (2007), Goodman & Weare (2010, hereafter GW) proposed an affine invariant sampling algorithm (§(2)) with only two hyperparameters that can be tuned for performance. Hou *et al.* (2011) were the first group to implement this algorithm to solve a physics problem. The implementation presented here is an independent effort that has already proved effective in several projects (Lang & Hogg 2011; Bovy *et al.* 2011; Dorman *et al.* 2012, Foreman-Mackey & Widrow 2012, in prep.). In what follows, we summarize the GW algorithm and the implementation decisions made in `emcee`. In §(2.2) we describe the small changes that must be made to the algorithm to parallelize it. Finally, in §(4.3), we outline the installation, usage and troubleshooting of the package.

2. The Algorithm

2.1. The stretch move

Goodman & Weare (2010) proposed an affine invariant ensemble sampling algorithm

informally called the “stretch move”. For completeness and for clarity of notation, we summarize the algorithm here and refer the interested reader to the original paper for more details. This method involves simultaneously evolving an ensemble of K *walkers* $\mathbf{X} = \{X_j, \forall j = 1, \dots, K\}$ where the proposal distribution for one walker k is based on the current positions of the $K - 1$ walkers in the *complementary ensemble* $\mathbf{X}_{[k]} = \{X_j, \forall j \neq k\}$. In general, each X_j is also a vector in n dimensions, where n is the dimension of the parameter space.

To update the position of a walker at position \mathbf{X}_k , another walker \mathbf{X}_j with $j \neq k$ is randomly chosen and then a new position is proposed:

$$X_k(t) \rightarrow Y = X_j + Z [X_k(t) - X_j] \quad (5)$$

where Z is a random variable drawn from a distribution $g(z)$. It is clear that if $g(z)$ satisfies

$$g(z^{-1}) = z g(z), \quad (6)$$

the proposal of Equation (5) is symmetric. In this case, the chain will satisfy detailed balance if the proposal is accepted with probability

$$q = \min \left\{ 1, Z^{n-1} \frac{p(\mathbf{Y})}{p(\mathbf{X}_k(t))} \right\} \quad (7)$$

where n is the dimension of the parameter space. This procedure is then repeated for each walker in the ensemble *in series* following the procedure shown in Algorithm (1).

[Goodman & Weare \(2010\)](#) advocate for a particular form of $g(z)$, namely

$$g(z) \propto \begin{cases} \frac{1}{\sqrt{z}} & \text{if } z \in \left[\frac{1}{a}, a \right], \\ 0 & \text{otherwise} \end{cases} \quad (8)$$

where a is an adjustable scale parameter that [Goodman & Weare \(2010\)](#) set to 2.

2.2. The parallel stretch move

It is tempting to naïvely parallelize the stretch move algorithm by simultaneously advancing each walker based on the state of the ensemble instead of evolving the walkers in series. Unfortunately, this would no longer satisfy detailed balance. Instead, if we split the ensemble into two ensembles ($\mathbf{X}^{(0)}$ and $\mathbf{X}^{(1)}$) and simultaneously update all the blue walkers — using the stretch move procedure — based on the positions of *only the red walkers* then

Algorithm 1 A single stretch move update step from Goodman & Weare (2010) where line 5 is generally the most computationally expensive step.

```

1: for  $k = 1, \dots, K$  do
2:   Draw a walker  $X_j$  at random from the complementary ensemble  $\mathbf{X}_{[k]}(t)$ 
3:    $Z \leftarrow z \sim g(z)$ , Equation (8)
4:    $Y \leftarrow X_j + Z [X_k(t) - X_j]$ 
5:    $q \leftarrow Z^{n-1} p(Y)/p(X_k(t))$ 
6:    $R \leftarrow r \sim [0, 1]$ 
7:   if  $R \geq q$ , Equation (7) then
8:      $X_k(t+1) \leftarrow Y$ 
9:   else
10:     $X_k(t+1) \leftarrow X_k(t)$ 
11:   end if
12: end for

```

the outcome is a valid step for each of the walkers. Then, the red walkers are advanced based only on the positions in the blue ensemble. The pseudocode for this procedure is shown in Algorithm (2). This code looks very similar to that in Algorithm (1), however, now the computationally expensive inner loop (starting at line 2 in Algorithm (2)) can be run in parallel.

The performance of this method — quantified by the autocorrelation time — is comparable to the traditional stretch move algorithm but the fact that one can now take advantage of generic parallelization makes this generalization extremely powerful.

3. Benchmarks & Tests

3.1. Measuring the performance

A standard method of quantifying the performance of an MCMC sampler is to estimate the autocorrelation time of the sampler on several densities.

The main goal of running a Markov chain is to measure the expectation value (and variance) of a particular value (e.g. f)

$$\langle f(\mathbf{x}) \rangle = \int f(\mathbf{x}) p(\mathbf{x}) d\mathbf{x} \quad (9)$$

Algorithm 2 The parallel stretch move update step

```

1: for  $i \in \{0, 1\}$  do
2:   for  $k = 1, \dots, K/2$  do
3:     Draw a walker  $X_j$  at random from the complementary ensemble  $\mathbf{X}^{(\sim i)}(t)$ 
4:      $Z \leftarrow z \sim g(z)$ , Equation (8)
5:      $Y \leftarrow X_j + Z [X_k^{(i)}(t) - X_j]$ 
6:      $q \leftarrow Z^{n-1} p(Y) / p(X_k^{(i)}(t))$ 
7:      $R \leftarrow r \sim [0, 1]$ 
8:     if  $R \geq q$ , Equation (7) then
9:        $X_k^{(i)}(t + \frac{1}{2}) \leftarrow Y$ 
10:    else
11:       $X_k^{(i)}(t + \frac{1}{2}) \leftarrow X_k(t)$ 
12:    end if
13:  end for
14:   $t \leftarrow t + \frac{1}{2}$ 
15: end for

```

which can be approximated as

$$\langle f(\mathbf{x}) \rangle \approx \frac{1}{T_s} \sum_{t=1}^{T_s} f(\mathbf{X}(t)) \quad (10)$$

where T is the length of the chain. The generalization of Equation (10) to the case of the ensemble sampler is

$$\langle f(\mathbf{x}) \rangle \approx \frac{1}{T_e} \sum_{t=1}^{T_s} \left[\frac{1}{K} \sum_{k=1}^K f(\mathbf{X}_k(t)) \right] \quad (11)$$

where K is the number of walkers. The autocorrelation function of the chain is then given by

$$C(t) = \frac{1}{K^2} \lim_{t' \rightarrow \infty} \text{cov} \left[\sum_{k=1}^K f(\mathbf{X}_k(t + t')), \sum_{k=1}^K f(\mathbf{X}_k(t')) \right] \quad (12)$$

and the integrated autocorrelation time is given by

$$\tau = \sum_{t=-\infty}^{\infty} \frac{C(t)}{C(0)}. \quad (13)$$

emcee depends on the Python module `acor`¹ to estimate the autocorrelation time. This

¹<http://github.com/dfm/acor>

module is a direct port of the original algorithm (described by [Goodman & Weare 2010](#)) and implemented by those authors in C++.²

3.2. Multivariate Gaussian distribution

The simplest test of an MCMC sampler is its sampling performance on a highly covariant multivariate Gaussian density. For the tests in the paper, we randomly generated a 50 dimensional positive definite covariance tensor and initial conditions for each walker. Then, each sampler was tested with the same initial conditions for its performance on the density

$$\pi(\mathbf{x}) \propto \exp\left(-\frac{1}{2}\mathbf{x}^T \Sigma^{-1} \mathbf{x}\right) \quad (14)$$

for the same covariance tensor Σ in each trial.

DFM: actually do some tests here...

4. Usage

4.1. Installation

The easiest way to install `emcee` is using `pip`³. Running the command

```
% pip install emcee
```

at the command line of a UNIX-based system will install the package and its `Python` dependencies. If you would like to install for all users, you might need to run the above command with superuser permissions. `emcee` depends on `Python` (> 2.7) and `numpy`⁴ (> 1.6) and the associated `dev` headers. On some systems, you might need to install these packages separately. On `Ubuntu`, you can install these dependencies using the command:

```
% apt-get python python-dev numpy numpy-dev
```

An alternative installation method is to download the source code from <http://danfm.ca/emcee> and run

²<http://www.math.nyu.edu/faculty/goodman/software/acor>

³<http://pypi.python.org>

⁴<http://numpy.scipy.org>

```
% python setup.py install
```

in that directory. Make sure that you have `numpy` installed as well.

4.2. Basic usage

```
import numpy as np
import emcee

def lnprobfn(x, mu, icov):
    diff = x-mu
    return -np.dot(diff,np.dot(icov,diff))/2.0

ndim = 10
means = np.random.rand(ndim)
cov = 0.5*np.random.rand(ndim**2).reshape((ndim, ndim))
cov = np.triu(cov)
cov += cov.T - np.diag(cov.diagonal())
cov = np.dot(cov,cov)
icov = np.linalg.inv(cov)

nwalkers = 100
p0 = [np.random.rand(ndim) for i in xrange(nwalkers)]

sampler = emcee.EnsembleSampler(nwalkers, ndim, lnprobfn,
                               args=[means, icov])

pos,prob,state = sampler.run_mcmc(p0, None, 500)
sampler.clear_chain()

sampler.run_mcmc(pos, state, 2000)

print "Mean acceptance fraction:", np.mean(sampler.acceptance_fraction)
```

4.3. API

4.3.1. *emcee.ensemble*

```
emcee.EnsembleSampler(*args, **kwargs)
```

Ensemble sampling following Goodman & Weare (2010) with optional parallelization

Arguments

`k (int)`: The number of Goodman & Weare “walkers”.

`dim (int)`: Number of dimensions in the parameter space.

`lnpostfn (callable)`: A function that takes a vector in the parameter space as input and returns the natural logarithm of the posterior probability for that position.

Keyword Arguments

`a (float)`: The proposal scale parameter from Equation (8). (default: 2.0)

`args (list)`: Optional list of extra arguments for `lnpostfn`. `lnpostfn` will be called with the sequence `lnpostfn(p, *args)`.

`postargs (list)`: Alias of `args` for backwards compatibility.

`threads (int)`: The number of threads to use for parallelization. If `threads == 1`, then the `multiprocessing` is not used but if `threads > 1`, then a `Pool` object is created and calls to `lnpostfn` are run in parallel following Algorithm (2). (default: 1)

`pool (multiprocessing.Pool)`: An alternative method of using the parallelized algorithm. If `pool` is not `None`, the value of `threads` is ignored and the provided `Pool` is used for all parallelization. (default: `None`)

Exceptions

`AssertionError`: If `k < 2*dim` or if `k` is not even.

Properties

`chain (numpy.ndarray)`: A pointer to the Markov chain itself. The shape of this array is `(k, dim, iterations/resample)`

`flatchain (numpy.ndarray)`: A shortcut for accessing `chain` flattened along the zeroth (walker) axis.

`lnprobability (numpy.ndarray)`: A pointer to the matrix of the value of `lnprobfn` produced at each step for each walker. The shape is `(k, iterations/resample)`.

`iterations (int)`: The number of steps that have been run in the chain.

`acceptance_fraction` (`numpy.ndarray`): An array (length: `k`) of the fraction of steps accepted for each walker.

`acor` (`numpy.ndarray`): The autocorrelation time of each parameter in the chain (length: `dim`) as estimated by the `acor` module.

`random_state` (`tuple`): The state of the internal random number generator. In practice, it's the result of calling `get_state()` on a `numpy.random.mtrand.RandomState` object. You can try to set this property but be warned that if you do this and it fails, it will do so silently.

Methods

`emcee.EnsembleSampler.sample(pos0, lnprob0=None, rstate0=None, iterations=1)`

Advances the chain `iterations` steps as an iterator

Arguments

`pos0` (`numpy.ndarray`): A list of the initial positions of the walkers in the parameter space. The shape is `(k, dim)`.

Keyword Arguments

`lnprob0` (`numpy.ndarray`): The list of log posterior probabilities for the walkers at positions given by the `p0`. If `lnprob` is `None`, the initial values are calculated. The shape is `(k, dim)`.

`rstate0` (`tuple`): The state of the random number generator.
See `EnsembleSampler.random_state` for details.

`iterations` (`int`): The number of steps to run. (default: 1)

Yields

`pos` (`numpy.ndarray`): A list of the current positions of the walkers in the parameter space. The shape is `(k, dim)`.

`lnprob` (`numpy.ndarray`): The list of log posterior probabilities for the walkers at positions given by the `pos`. The shape is `(k, dim)`.

`rstate` (`tuple`): The state of the random number generator.

`emcee.EnsembleSampler.run_mcmc(pos0, N, rstate0=None, lnprob0=None, **kwargs)`

Iterate `sample` for `N` iterations and return the result. The parameters are passed directly to `sample` so see above for details.

Returns

`pos` (`numpy.ndarray`): A list of the final positions of the walkers in the parameter space. The shape is `(k, dim)`.

`lnprob` (`numpy.ndarray`): The list of log posterior probabilities for the walkers at positions given by the `pos`. The shape is `(k, dim)`.

`rstate` (`tuple`): The final state of the random number generator.

```
emcee.EnsembleSampler.reset()
```

Clear `chain`, `lnprobability` and the bookkeeping parameters.

```
emcee.EnsembleSampler.clear_chain()
```

An alias of `reset` kept for backwards compatibility.

4.4. Issues & Contributions

The development of `emcee` is being coordinated on GitHub at <http://github.com/dfm/emcee> and contributions are welcome. If you encounter any problems with the code, please report them at <http://github.com/dfm/emcee/issues> (DFM: check this url) and consider contributing a patch.

REFERENCES

- Bovy, J., Rix, H.-W., Liu, C., Hogg, D. W., Beers, T. C., & Lee, Y. S., 2011, ApJ, submitted, arXiv:1111.1724 [astro-ph.GA]
- Christen, J., *A general purpose scale-independent MCMC algorithm*, technical report I-07-16, CIMAT, Guanajuato, 2007.
- Dorman, C., Guhathakurta, P., Fardal, M. A., Geha, M. C., Howley, K. M., Kalirai, J. S., Lang, D., Cuillandre, J., Dalcanton, J., Gilbert, K. M., Seth, A. C., Williams, B. F., & Yniguez, B., 2012, ApJ, submitted

- Dunkley, J., Bucher, M., Ferreira, P. G., Moodley, K., & Skordis, C., 2005, MNRAS, 356, 925-936
- Goodman, J., & Weare, J., 2010, Comm. App. Math. Comp. Sci., 5, 65
- Gregory, P. C., *Bayesian Logical Data Analysis for the Physical Sciences*, Cambridge University Press, 2005
- Hou, F., Goodman, J., Hogg, D. W., Weare, J., & Schwab, C., 2011, arXiv:1104.2612
- Lang, D. and Hogg, D. W., 2011, arXiv:1103.6038
- MacKay, D., *Information Theory, Inference, and Learning Algorithms*, Cambridge University Press, 2003
- Widrow, L. M. and Pym, B. and Dubinski, J., 2008, ApJ, 679, 1239